

MPTA-Repair: Simultaneous Multi-Property Clock Bound Repair for Networks of Timed Automata via Iterative Space Optimization

Guangtong Zhou, Run Chang, Ji Wu*

School of Computer Science and Engineering, Beihang University, Beijing, China

Email: zhougt@buaa.edu.cn, handsomeruncr@gmail.com, wuji@buaa.edu.cn

Abstract—Networks of timed automata (NTA) are widely used to model and verify real-time industrial systems, which are typically required to satisfy multiple properties simultaneously. When verification fails, repairing an NTA is challenging because a valid repair must modify erroneous clock constraints, satisfy all properties, and preserve the original functional behavior. We present MPTA-Repair, an automated repair framework for multi-property NTA. It uses an iterative, counterexample-guided MaxSMT approach to generate minimally modified candidates, exploiting relations among properties, counterexamples, and repair variables to optimize search. Each candidate is validated by full-property regression verification and an admissibility check based on functional equivalence. Compared to the TARTAR baseline, MPTA-Repair achieves higher effectiveness on faulty models generated via our proposed mutation strategy. Specifically, it improves the success rate from 72% to 93% on benchmarks derived from UPPAAL and the XtaBenchmarkSuite, and from 20% to 76% on an industrial dataset constructed in this work.

Index Terms—networks of timed automata, automated model repair, MaxSMT, multi-property verification

I. INTRODUCTION

Real-time industrial systems combine discrete control logic with quantitative timing requirements. In railway control, communication protocols, and aerospace scheduling [1]–[3], correctness depends not only on which actions occur, but also on whether they occur within permitted time windows. The scope of this paper is timed safety properties rather than unbounded liveness properties, because deadlines, bounded waiting, alarm durations, and error-state reachability describe finite executions whose violations are directly observable during design and testing [4], [5]. Timed automata (TAs) model such behavior with clocks, guards, resets, and invariants [6], [7]; practical systems are often modeled as networks of timed automata (NTA).

Clock behavior is central to timed safety verification. Clocks and their constraints encode deadlines, timeouts, durations, and delays. An incorrect numerical bound in a transition guard or location invariant can make otherwise valid discrete behavior occur too early, too late, or for an invalid duration [2], [4], [8]. This paper studies clock-bound repair for NTA models, modifying only numerical bounds in guards and invariants while leaving properties, locations, transitions,

synchronizations, clock references, resets, and data updates unchanged. Model checkers such as UPPAAL, Kronos, and opaal can return timed diagnostic traces (TDTs), namely timed counterexamples reaching violating states [9]–[11]. A TDT exposes a violating execution, but it does not identify a globally acceptable clock-bound edit.

Repair is more difficult when an NTA must satisfy multiple properties simultaneously. A bound change that blocks one TDT may leave another violation feasible, violate a previously satisfied property, or compromise the original functional behavior. In this paper, functional behavior means the untimed discrete behavior characterized by functional equivalence. Thus, local trace elimination is necessary but insufficient: a candidate edit should be accepted only when it satisfies the full property set and preserves the original functional behavior. A practical repair method should therefore generate small and inspectable candidate edits from local TDT evidence, but accept a candidate only after full-property regression verification and an admissibility check based on untimed functional equivalence [4].

Existing repair techniques provide an important foundation. TARTAR [4] introduced TDT-based clock-bound repair by encoding a TDT into linear real arithmetic constraints, solving a MaxSMT problem for small bound changes, and checking admissibility through functional equivalence. Related work studies parametric timed automata, clock-guard repair, robustness analysis, and SMT-based repair [12]–[15]. However, these workflows are primarily trace-local and single-property in style. When repairs are generated property by property, the coupling among properties, counterexamples, and repairable bounds is not used to guide candidate generation, and a local edit may be rejected only after full validation.

This paper presents MPTA-Repair, an automated framework for simultaneous multi-property clock-bound repair of NTA models via iterative solving. The key idea is to treat each violated property and its TDT as a property-trace constraint, maintain a shared constraint pool, and use relations among properties, counterexamples, and repairable bounds to guide MaxSMT-based candidate generation. Each candidate is validated by full-property regression verification and an admissibility check based on untimed functional equivalence. If validation exposes new violations, their TDTs are added to the pool and the solver is invoked again.

The implementation and datasets are publicly available at <https://github.com/zhougt/MPTA-Repair>.

We evaluate MPTA-Repair on benchmarks derived from UPPAAL [16] and the XtaBenchmarkSuite [17], together with an industrial dataset constructed in this work. Faulty NTA instances are produced by controlled clock-bound mutations. Compared with an adapted iterative TARTAR-style baseline, MPTA-Repair improves the success rate from 72% to 93% on the benchmark dataset and from 20% to 76% on the industrial dataset.

This paper makes three contributions. First, it formulates simultaneous multi-property clock-bound repair for NTA models under full-property satisfaction and preservation of untimed functional behavior. Second, it introduces an iterative MaxSMT workflow based on a shared property-trace constraint pool and relation-guided candidate generation. Third, it evaluates the framework on benchmark and industrial faulty models, showing that joint multi-property repair is more effective than an iterative trace-local baseline in this setting.

The remainder of the paper is organized as follows. Section II introduces the notation and accepted-repair condition. Section III presents MPTA-Repair. Section IV reports the experimental evaluation. Section V concludes the paper.

II. PRELIMINARIES

A. Networks of Timed Automata and Properties

The paper follows standard timed-automata semantics [7]. This section establishes the notation used for multi-property clock-bound repair.

Definition 1 (Timed automaton). A timed automaton (TA) is a tuple

$$T = (L, \ell_0, C, \Sigma, \Theta, I), \quad (1)$$

where L is a finite set of locations, ℓ_0 is the initial location, C is a finite set of clocks, Σ is a set of action labels, Θ is a finite set of transitions, and I maps each location to a clock invariant. A clock constraint is a conjunction of atoms $x \sim b$, where $x \in C$, $\sim \in \{<, \leq, =, \geq, >\}$, and $b \in \mathbb{R}_{\geq 0}$. A transition $\theta = (\ell, g, a, R, \ell')$ moves from ℓ to ℓ' when guard g is satisfied, emits action a , and resets clocks in $R \subseteq C$. Each numerical constant b occurring in a guard or invariant is a clock-bound occurrence. MPTA-Repair only changes repairable clock-bound occurrences; it does not modify locations, transitions, synchronization labels, clock references, resets, data updates, or properties.

Definition 2 (NTA with property set). An NTA specification consists of a network of timed automata and the properties that the network must satisfy:

$$\begin{aligned} M &= (N, \Phi), \\ N &= T_1 \parallel \dots \parallel T_k, \\ \Phi &= \{\varphi_1, \dots, \varphi_m\}. \end{aligned} \quad (2)$$

The network N models a real-time system by composing interacting timed automata, with synchronization semantics supplied by the target modeling language, such as UPPAAL [9]. The property set Φ is part of the system specification: it states the timing and behavioral requirements that the NTA must

satisfy. For convenience, $N \models \Phi$ denotes that every property in Φ holds over the reachable states of N .

The repair problem considered here is restricted to timed safety properties, which require a state predicate to hold over all reachable states. In UPPAAL-style notation, such properties often have the invariant form $A[] p$, where p may combine location predicates, data predicates, and clock constraints. Clocks are central to timed safety properties because deadlines, timeouts, durations, and earliest or latest event times are expressed by numerical bounds over clock values.

B. Timed Diagnostic Traces and Accepted Repairs

Definition 3 (Timed diagnostic trace). Given an NTA specification $M = (N, \Phi)$ and a property $\varphi \in \Phi$, a timed diagnostic trace (TDT) for φ is a feasible finite timed execution of N that reaches a state violating φ . In practice, a TDT is returned by a model checker when verification of φ fails. It records the relevant locations and transitions, guards and invariants, clock resets, timing information such as delays or zones, and the violation point.

In MPTA-Repair, each TDT is treated as a property-trace constraint: the violated property identifies the requirement to be restored, and the trace identifies the timed behavior that makes the violation feasible. Section III uses such constraints from multiple properties in a joint solving process.

Definition 4 (Untimed functional equivalence). For an NTA N , let $L_\mu(N)$ denote the untimed language obtained by abstracting from delays while preserving the discrete action or transition behavior feasible under the timing constraints. Two NTA models N_1 and N_2 are untimed functionally equivalent if

$$L_\mu(N_1) = L_\mu(N_2). \quad (3)$$

This notion follows the functional-equivalence criterion used for clock-bound repair [4]: a timing repair should not remove intended discrete behavior or add new discrete behavior, even though it changes timing constraints to eliminate violations.

Definition 5 (Accepted repair). Let $M_{bad} = (N_{bad}, \Phi)$ be a faulty NTA specification. A candidate repaired model $M_{rep} = (N_{rep}, \Phi)$ is obtained by changing only repairable clock-bound occurrences in guards and invariants. The candidate is accepted as a repair only if

$$N_{rep} \models \Phi \quad \wedge \quad L_\mu(N_{bad}) = L_\mu(N_{rep}). \quad (4)$$

The first conjunct corresponds to full-property regression verification: the repaired model must satisfy the entire property set, not only the property that produced the current TDT. The second conjunct is the admissibility check based on untimed functional equivalence; it prevents repairs that make properties true by disabling required functional behavior. MPTA-Repair uses this accepted-repair condition as the global validation criterion for candidates generated from local TDT evidence.

III. APPROACH

A. Overview

Let N be an NTA and $\Phi = \{\varphi_1, \dots, \varphi_m\}$ a set of timed safety properties. We write $M = (N, \Phi)$ for the repair input.

MPTA-Repair changes only repairable numerical clock bounds in guards and invariants. It returns a repaired specification $M' = (N', \Phi)$ only when N' satisfies the accepted-repair condition in Eq. (4); otherwise, it reports that no repair was found within the resource budget.

The main difficulty is that the properties in Φ do not define independent repair tasks. Different properties may constrain the same clock bound, overlapping execution paths, or different stages of a shared timing budget. Their TDTs may also have different relations: some are duplicates, some share only part of a discrete path, and some reach different violations through the same repairable bounds. Consequently, an edit derived from one TDT may leave another violating execution feasible or invalidate a property that already holds. A single-property workflow that processes these TDTs independently does not retain their joint constraints and may repeatedly explore incompatible candidate edits.

MPTA-Repair uses these relations to organize the current repair problem before MaxSMT solving. At the property level, equivalent or stronger requirements are identified to avoid redundant constraints. At the trace level, duplicate and structurally related TDTs are grouped. At the repair-bound level, each property-trace constraint is mapped to the bounds that can affect it. Property-trace constraints that depend on disjoint bound sets form separate search groups, whereas constraints that share repairable bounds remain in the same group and are solved jointly. This partition reduces independent parts of the search without separating constraints that require coordinated edits.

A trace-local repair workflow encodes one TDT, solves a MaxSMT instance, and validates the resulting candidate. MPTA-Repair extends this workflow with a shared property-trace pool, relation-guided selection of active constraints and bounds, and iterative refinement across properties. At refinement round k , $\mathcal{T}_k$ denotes the current set of property-trace pairs (τ, φ) , where τ is a TDT for $\varphi \in \Phi$. The framework verifies N' against the full property set, updates $\mathcal{T}_k$, constructs a joint MaxSMT problem for the dependent groups, and instantiates the resulting clock-bound edits in a candidate model.

A candidate is not accepted merely because it satisfies the active MaxSMT constraints. It must first pass full-property regression verification and then the admissibility check based on functional equivalence. If either check fails, newly exposed TDTs are added to the pool and the rejected assignment is excluded from later rounds. Only a candidate that passes both checks is returned as M' . Figure 1 summarizes this iterative workflow.

B. Joint Trace Encoding and Optimization Objective

MPTA-Repair follows the clock-bound variation strategy used in TARTAR [4]. Let S denote the set of repairable clock-bound occurrences. For each occurrence $s \in S$, the original numerical bound b_s in a transition guard or location invariant is associated with a variation variable δ_s . This ensures that

the repaired constraint retains the original clock reference and comparison operator while incorporating a shifted value:

$$b'_s = b_s + \delta_s \quad (5)$$

The repair space is therefore strictly restricted to clock-bound constants, keeping transition structures and discrete updates unaltered.

MPTA-Repair builds on the trace-encoding formulation of TARTAR and extends it to the multi-property setting [4]. Given a TDT

$$\tau = \ell_0 \xrightarrow{d_0, \theta_0} \ell_1 \cdots \xrightarrow{d_{n-1}, \theta_{n-1}} \ell_n \quad (6)$$

where n is the number of discrete transitions, d_i is the delay before the i th transition, and $\theta_i = (\ell_i, g_i, a_i, R_i, \ell_{i+1})$ is the i th transition. Let $\rho = (b'_s)_{s \in S}$ denote the modified-bound vector, and let $\mathbf{z}_\tau = (d_0, \dots, d_{n-1}, \nu_0, \dots, \nu_n)$ collect the delays and clock valuations, where ν_i is the clock valuation at location ℓ_i . Then the trace path is encoded as a conjunction of linear arithmetic constraints:

$$\begin{aligned} \text{Path}(\tau, \rho, \mathbf{z}_\tau) = & \text{Init}(\nu_0) \\ & \wedge \bigwedge_{i=0}^{n-1} \left(d_i \geq 0 \wedge \text{Inv}(\ell_i, \nu_i + d_i, \rho) \right. \\ & \wedge \text{Guard}(\theta_i, \nu_i + d_i, \rho) \\ & \left. \wedge \nu_{i+1} = \text{Reset}(\nu_i + d_i, R_i) \right) \end{aligned} \quad (7)$$

Here, $\text{Init}(\nu_0)$ constrains the initial clock valuation, Inv evaluates the invariant $I(\ell_i)$, and Guard evaluates the guard g_i of transition θ_i under the given valuation and modified bounds. The reset equation $\nu_{i+1} = \text{Reset}(\nu_i + d_i, R_i)$ updates the valuation after transition θ_i by resetting clocks in R_i and preserving the others. The original violation encoding combines path executability with the final property violation:

$$\text{Enc}(\tau, \varphi, \rho) = \text{Path}(\tau, \rho, \mathbf{z}_\tau) \wedge \text{Bad}_\varphi(\ell_n, \nu_n) \quad (8)$$

The predicate $\text{Bad}_\varphi(\ell_n, \nu_n)$ characterizes the violation state of φ at the final location. To block this violating behavior while preserving executability of the discrete path skeleton, the single-trace repair constraint uses two disjoint copies $\mathbf{z}_\tau^+$ and $\mathbf{z}_\tau^-$ of the delay and clock-valuation variables:

$$\begin{aligned} \Psi_{\tau, \varphi}^{\text{repair}}(\rho) = & \exists \mathbf{z}_\tau^+. \text{Path}(\tau, \rho, \mathbf{z}_\tau^+) \\ & \wedge \neg \exists \mathbf{z}_\tau^-. \left(\text{Path}(\tau, \rho, \mathbf{z}_\tau^-) \right. \\ & \left. \wedge \text{Bad}_\varphi(\ell_n, \nu_n^-) \right) \end{aligned} \quad (9)$$

The first part requires that the original discrete path skeleton still has at least one timed realization after repair. The second part requires that no assignment of delays and clock valuations can make that same path violate φ .

We extend this single-trace formulation to a joint, multi-property incremental setting. At refinement round k , counterexamples from all currently violated properties are accumulated into a trace pool

$$\mathcal{T}_k = \{(\tau_j, \varphi_j)\}_{j=1}^{m_k} \quad (10)$$

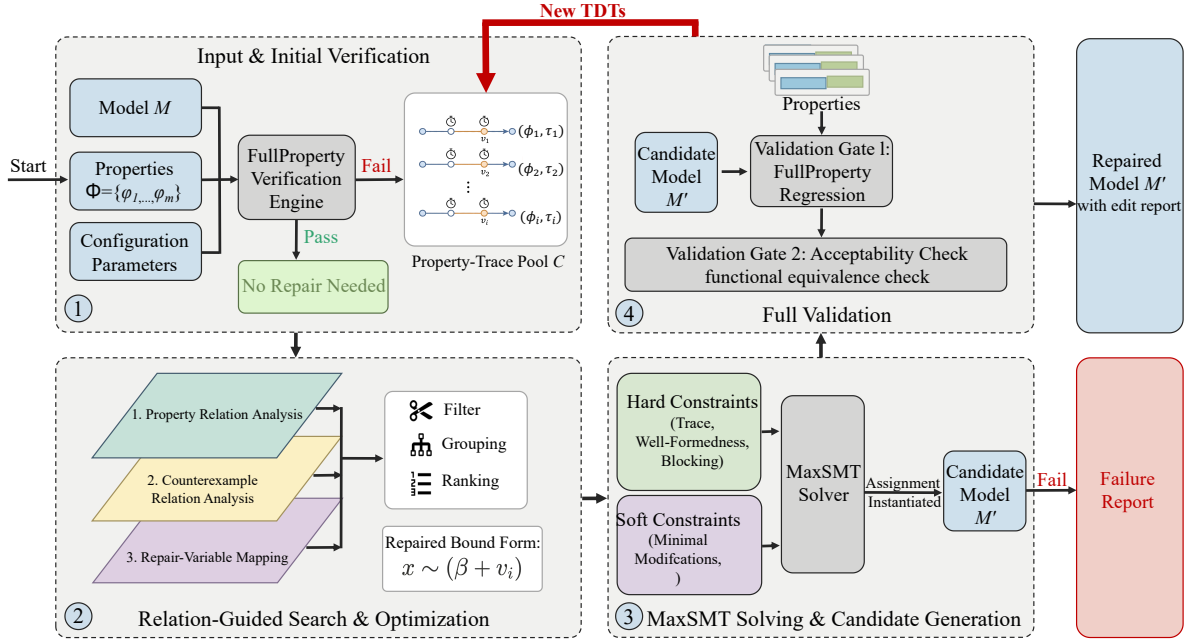


Fig. 1. Overview of the MPTA-Repair workflow.

where $m_k = |\mathcal{T}_k|$ is the number of property-trace pairs in the pool at round k . MPTA-Repair assembles a joint MaxSMT problem with the constraint system defined as

$$\mathcal{F}_k(\rho) = \text{Dom}(\rho) \wedge \text{Block}_k(\rho) \wedge \bigwedge_{(\tau, \varphi) \in \text{Select}(\mathcal{T}_k)} \Psi_{\tau, \varphi}^{\text{repair}}(\rho) \quad (11)$$

Here, $\text{Dom}(\rho)$ enforces static interval consistency and domain limits, such as $b'_s \geq 0$; $\text{Block}_k(\rho)$ collects the blocking clauses accumulated before round k and is initially true; and $\text{Select}(\mathcal{T}_k)$ represents a reduced subset of traces obtained after relation-aware filtering. When an instantiated candidate NTA fails verification, the constraint system is updated incrementally. The framework expands the pool via

$$\mathcal{T}_{k+1} = \mathcal{T}_k \cup \text{NewTDTs}(N(\rho_k), \Phi) \quad (12)$$

where ρ_k is the modified-bound assignment produced in round k , $N(\rho_k)$ is the candidate network obtained by applying ρ_k to N , and $\text{NewTDTs}(N(\rho_k), \Phi)$ is the set of newly exposed property-trace pairs returned by full-property verification. The framework then updates the search-space block to exclude the failed assignment ρ_k over the set of repairable bounds S :

$$\text{Block}_{k+1}(\rho) = \text{Block}_k(\rho) \wedge \left(\bigvee_{s \in S} \delta_s \neq \delta_s^k \right) \quad (13)$$

where δ_s^k is the valuation assigned to variation variable δ_s in round k .

The optimization objective is governed by a hierarchical set of soft constraints. MPTA-Repair uses a lexicographic

objective that first favors small repairs and then incorporates risk- and coverage-based priorities:

$$\text{lex min}_{\rho} \left(\begin{array}{l} \sum_{s \in S} \mathbf{1}\{\delta_s \neq 0\}, \\ \sum_{s \in S} |\delta_s|, \\ \sum_{s \in S} \text{risk}_s \cdot \mathbf{1}\{\delta_s \neq 0\}, \\ - \sum_{s \in S} \text{benefit}_s \cdot \mathbf{1}\{\delta_s \neq 0\} \end{array} \right) \quad (14)$$

Here, $\mathbf{1}\{\cdot\}$ denotes the indicator function, yielding 1 when the inner condition holds and 0 otherwise. The primary and secondary objectives minimize the number and total magnitude of altered clock bounds to produce concise repairs. The third term incorporates a penalty coefficient risk_s to discourage modifications prone to regression, while the final term uses a reward coefficient benefit_s to prioritize repair sites that cover a broader set of violated traces.

C. Relation-Guided Search Optimization

During refinement, accumulating violated properties, TDTs, and editable bounds expands the MaxSMT search space. To reduce redundant constraints, MPTA-Repair uses relation-guided optimization to filter, group, and rank constraints and repair variables. These heuristics only guide the search; correctness is still guaranteed through full-property regression verification and an admissibility check.

At the property layer, timed safety properties are normalized into an $\mathbf{A}[\]$ fragment over location predicates and clock or integer constraints. Equivalence is detected by the syntactic

TABLE I
STATISTICS OF THE BENCHMARK AND INDUSTRIAL DATASETS.

Dataset	Original models	Mutated models	Avg. properties/model
Benchmark	9	54	4.1
Industrial	5	105	7.3

identity of normalized forms, while dominance is evaluated under logical subsumption, e.g., $A[]x \leq 5$ subsumes $A[]x \leq 10$ within the same structural context. MPTA-Repair then retains only the stronger properties for solving to optimize efficiency.

At the counterexample layer, MPTA-Repair maintains a TDT pool indexed by violated properties. Exact duplicates are removed, and the remaining traces are annotated with traversed transitions, visited locations, active clocks, and guard or invariant bounds. Traces with identical discrete transition sequences or overlapping clock zones are grouped to reveal common bottlenecks. A trace constraint is omitted from the active encoding only when it is an exact duplicate or soundly subsumed by an encoded trace constraint.

At the repair-bound layer, MPTA-Repair maps each active property-trace constraint to the clock bounds occurring in or affecting its encoding. This mapping restricts active MaxSMT decision variables to bounds relevant to the current counterexample pool and ranks them for candidate generation. Bounds with higher relational scores are prioritized because their modifications are more likely to satisfy the joint constraint system.

IV. EXPERIMENTAL EVALUATION

This section evaluates MPTA-Repair on both benchmark and industrial case studies. The evaluation follows the clock-bound repair setting defined in Section III, where only numerical bounds in guards and invariants are repairable. For comparison, we evaluate MPTA-Repair against an adapted iterative TARTAR-style baseline [4].

A. Evaluation Setup and Research Questions

We evaluate MPTA-Repair on two datasets, summarized in Table I. The benchmark dataset includes nine benchmark models derived from UPPAAL benchmark [16] and the XtaBenchmarkSuite [17], such as Elevator and FDDI. The industrial dataset constructed in this work contains five closely coupled NTA case-study models: Gear Control System for automotive transmission [18], SAE RoR for autonomous-driving decisions [19], Train Controller Alarm for railway emergency timing [20], Train Gate Controller for shared-resource concurrency [6], and Pacemaker for medical pacing [8]. Compared with the benchmark models, these case studies involve more tightly coupled timing constraints across multiple properties.

Since real-world NTA models with labeled clock-bound faults are difficult to obtain, we synthesize faulty instances using mutation [21], retaining only mutants that are syntactically valid, violate at least one timing safety property, and preserve the original functional behavior. Each repair instance

TABLE II
REPAIR RESULTS ACROSS BENCHMARK AND INDUSTRIAL DATASETS.

Dataset	Method	Mutated models	Repaired	Success rate	Avg. time
Benchmark	Iterative baseline	54	39	72%	60.61s
	MPTA-Repair	54	50	93%	15.46s
Industrial	Iterative baseline	105	21	20%	35.59s
	MPTA-Repair	105	80	76%	85.28s

is executed with a 10-minute timeout for both MPTA-Repair and the baseline.

We compare MPTA-Repair against an adapted iterative TARTAR-style baseline [4], which was originally designed for single-property repair and operates via trace-local repair generation. The evaluation is organized around the following research questions:

RQ1. How effective is MPTA-Repair compared with the iterative TARTAR-style baseline?

RQ2. How does repair effectiveness change from benchmark models to industrial NTA models with more tightly coupled timing constraints?

RQ3. What practical lessons do the results reveal about full-property validation, admissibility checking, and remaining repair failures?

B. Results by Research Question

RQ1: Overall repair effectiveness. Table II shows that MPTA-Repair outperforms the iterative baseline on both datasets. On the benchmark dataset, the baseline repairs 39 of 54 faulty instances, achieving a 72% success rate, whereas MPTA-Repair repairs 50 instances and reaches 93%. This result indicates that joint property-trace candidate generation improves repair effectiveness even when injected timing faults are relatively localized.

RQ2: Effect on industrial NTA models. The advantage of MPTA-Repair becomes more pronounced on the industrial dataset. The iterative baseline repairs only 21 of 105 faulty instances, corresponding to a 20% success rate, while MPTA-Repair repairs 80 instances and reaches 76%. The baseline reports a lower average duration on the industrial dataset because unsuccessful runs typically terminate early when no valid candidate can be found. Its success rate drops because trace-local repairs often produce candidates that fail full-property regression verification or the admissibility check under concurrent component interactions and tightly coupled timing constraints.

RQ3: Practical lessons and failure cases. MPTA-Repair achieves higher success rates by integrating constraints from TDTs and properties into a joint generation process. Consequently, the framework produces coordinated modifications that pass both full-property regression verification and the admissibility check. The results highlight two practical lessons for applying automated multi-property repair to NTA. First, full-property regression verification is necessary because industrial properties are often interdependent; a localized repair for one response-time property can inadvertently violate

another safety property, such as maximum waiting bounds. MPTA-Repair avoids these regressions by requiring each candidate to satisfy the full property set. Second, the admissibility check is critical to distinguish valid repairs from behavior-breaking edits that satisfy timing properties only by removing intended functional behavior, such as preventing critical system actions.

Nevertheless, the framework can still fail to generate repairs in certain cases when it fails to find an admissible candidate within the time budget. This limitation is prominent in complex models like the Pacemaker case study, where tightly coupled timing constraints expand the search space, causing the solver to time out before finding a valid candidate.

V. CONCLUSION

This paper presented MPTA-Repair, an automated workflow for multi-property and multi-counterexample clock-bound repair of industrial NTA models. To overcome the limitations of isolated single-trace fixes, the method accumulates TDTs and exploits relations among properties, counterexamples, and repairable bounds to guide MaxSMT-based search. Each candidate is validated through full-property regression verification and an admissibility check based on functional equivalence [4], ensuring that accepted repairs preserve the original functional behavior.

Evaluations across benchmark and industrial datasets show that MPTA-Repair outperforms the iterative baseline, particularly in industrial systems with dense property interactions. The failure analysis further indicates that unsuccessful repairs mainly arise when the framework cannot find an admissible candidate within the time budget, especially as dense timing interactions enlarge the clock-bound repair search space.

Future work will primarily focus on extending the automated repair space beyond numerical clock bounds to include broader structural modifications, such as altering synchronization labels and redesigning discrete control logic. We also plan to integrate the admissibility check earlier in repair generation, so that functionally invalid candidates can be pruned before costly validation and large-scale industrial repair can be made more efficient.

REFERENCES

- [1] D. Basile, A. Fantechi, and I. Rosadi, “Formal analysis of the UNISIG safety application intermediate sub-layer: Applying formal methods to railway standard interfaces,” in *Proceedings of Formal Methods for Industrial Critical Systems*, ser. Lecture Notes in Computer Science, vol. 12863. Springer, 2021, pp. 174–190.
- [2] H. Lönn and P. Pettersson, “Formal verification of a TDMA protocol start-up mechanism,” in *Proceedings of the 1997 IEEE Pacific Rim International Symposium on Fault-Tolerant Systems*, 1997, pp. 235–242.
- [3] M. Mikučionis, K. G. Larsen, J. I. Rasmussen, B. Nielsen, A. Skou, S. U. Palm, J. S. Pedersen, and P. Hougaard, “Schedulability analysis using Uppaal: Herschel-planck case study,” in *Proceedings of Leveraging Applications of Formal Methods, Verification and Validation*, ser. Lecture Notes in Computer Science, vol. 6416. Springer, 2010, pp. 175–190.
- [4] M. Kölbl, S. Leue, and T. Wies, “Clock bound repair for timed systems,” in *International Conference on Computer Aided Verification*. Springer, 2019, pp. 79–96.
- [5] T. Vogel, M. Carwehl, G. N. Rodrigues, and L. Grunske, “A property specification pattern catalog for real-time system verification with UPPAAL,” 2022, arXiv:2211.03817.
- [6] R. Alur and D. L. Dill, “A theory of timed automata,” *Theoretical Computer Science*, vol. 126, no. 2, pp. 183–235, 1994.
- [7] J. Bengtsson and W. Yi, “Timed automata: Semantics, algorithms and tools,” in *Lectures on Concurrency and Petri Nets*, ser. Lecture Notes in Computer Science. Springer, 2004, vol. 3098, pp. 87–124.
- [8] Z. Jiang, M. Pajic, A. Connolly, S. Dixit, and R. Mangharam, “Real-time heart model for implantable cardiac device validation and verification,” in *Proceedings of the Euromicro Conference on Real-Time Systems*. IEEE Computer Society, 2010, pp. 239–248.
- [9] K. G. Larsen, P. Pettersson, and W. Yi, “UPPAAL in a nutshell,” *International Journal on Software Tools for Technology Transfer*, vol. 1, no. 1–2, pp. 134–152, 1997.
- [10] C. Daws, A. Olivero, S. Tripakis, and S. Yovine, “The tool KRONOS,” in *Hybrid Systems III*, ser. Lecture Notes in Computer Science. Springer, 1996, vol. 1066, pp. 208–219.
- [11] A. E. Dalsgaard, R. R. Hansen, K. Y. Jørgensen, K. G. Larsen, M. C. Olesen, P. Olsen, and J. Srba, “opaal: A lattice model checker,” in *NASA Formal Methods*, ser. Lecture Notes in Computer Science, vol. 6617. Springer, 2011, pp. 487–493.
- [12] R. Alur, T. A. Henzinger, and M. Y. Vardi, “Parametric real-time reasoning,” in *Proceedings of the Twenty-Fifth Annual ACM Symposium on Theory of Computing*. ACM, 1993, pp. 592–601.
- [13] Étienne André, P. Arcaini, A. Gargantini, and M. Radavelli, “Repairing timed automata clock guards through abstraction and testing,” in *Proceedings of Tests and Proofs*, ser. Lecture Notes in Computer Science, vol. 11823. Springer, 2019, pp. 129–146.
- [14] J. Bendík, A. Sencan, E. A. Gol, and I. Černá, “Timed automata robustness analysis via model checking,” *Logical Methods in Computer Science*, vol. 18, no. 3, pp. 12:1–12:32, 2022.
- [15] M. Jose and R. Majumdar, “Bug-assist: Assisting fault localization in ANSI-C programs,” in *Proceedings of Computer Aided Verification*, ser. Lecture Notes in Computer Science, vol. 6806. Springer, 2011, pp. 504–509.
- [16] G. Behrmann, A. David, and K. G. Larsen, “A tutorial on Uppaal,” in *Formal Methods for the Design of Real-Time Systems*, ser. Lecture Notes in Computer Science, vol. 3185. Springer, 2004, pp. 200–236.
- [17] R. Farkas and G. Bergmann, “Towards reliable benchmarks of timed automata,” in *Proceedings of the 25th PhD Mini-Symposium*, 2018, pp. 20–23.
- [18] M. Lindahl, P. Pettersson, and W. Yi, “Formal design and analysis of a gear controller,” in *Proceedings of Tools and Algorithms for the Construction and Analysis of Systems*, ser. Lecture Notes in Computer Science, vol. 1384. Springer, 1998, pp. 281–297.
- [19] G. V. Alves, L. A. Dennis, and M. Fisher, “A double-level model checking approach for an agent-based autonomous vehicle and road junction regulations,” *Journal of Sensor and Actuator Networks*, vol. 10, no. 3, p. 41, 2021.
- [20] A. González, M. Cristiá, and C. Luna, “Verification of quantitative temporal properties in RealTime-DEVS,” *SIMULATION: Transactions of The Society for Modeling and Simulation International*, vol. 101, no. 10, pp. 1057–1076, 2025.
- [21] Y. Jia and M. Harman, “An analysis and survey of the development of mutation testing,” *IEEE Transactions on Software Engineering*, vol. 37, no. 5, pp. 649–678, 2011.