

MPTA-Repair: Simultaneous Multi-Property Clock Bound Repair for Networks of Timed Automata via Iterative Space Optimization

Guangtong Zhou, Run Chang, Ji Wu*

School of Computer Science and Engineering, Beihang University, Beijing, China

Email: zhougt@buaa.edu.cn, handsomeruncr@gmail.com, wuji@buaa.edu.cn

Abstract—Networks of timed automata (NTA) are widely used to model and verify real-time industrial systems, which are typically required to satisfy multiple properties simultaneously. When verification fails, repairing an NTA is challenging because a valid repair must modify erroneous clock constraints, satisfy all properties, and preserve the original functional behavior. We present MPTA-Repair, an automated repair framework for multi-property NTA. It uses an iterative, counterexample-guided MaxSMT approach to generate minimally modified candidates, exploiting relations among properties, counterexamples, and repair variables to optimize search. Each candidate is validated by full-property regression verification and an admissibility check based on functional equivalence. Compared to the TARTAR baseline, MPTA-Repair achieves higher effectiveness on faulty models generated via our proposed mutation strategy. Specifically, it improves the success rate from 72% to 93% on benchmarks derived from UPPAAL and the XtaBenchmarkSuite, and from 20% to 76% on an industrial dataset constructed in this work.

Index Terms—networks of timed automata, automated model repair, MaxSMT, multi-property verification

I. INTRODUCTION

Real-time industrial systems combine discrete control logic with quantitative timing requirements. In railway control, communication protocols, and aerospace scheduling [1]–[3], correctness depends not only on which actions occur, but also on whether they occur within permitted time windows. We focus on timed safety properties rather than unbounded liveness properties, because deadlines, bounded waiting, alarm durations, and error-state reachability describe finite executions whose violations are directly observable during design and testing [4], [5]. Timed automata (TAs) model such behavior with clocks, guards, resets, and invariants [6], [7]; practical systems are often modeled as networks of timed automata (NTA).

Within timed safety verification, clock behavior is central. Clocks and their constraints encode deadlines, timeouts, durations, and delays. An incorrect numerical bound in a transition guard or location invariant can make otherwise valid discrete behavior occur too early, too late, or for an invalid duration [2], [4], [8]. This paper therefore studies clock-bound repair for NTA models, modifying only numerical bounds in guards and invariants while leaving properties, locations, transitions,

synchronizations, clock references, resets, and data updates unchanged. Model checkers such as UPPAAL, Kronos, and opaal can return timed diagnostic traces (TDTs), i.e., timed counterexamples reaching violating states [9]–[11]. A TDT is useful failure evidence, but it is not a repair instruction: it exposes failures but not which bound is wrong or how to change it.

Repair is more difficult when an NTA must satisfy multiple properties simultaneously. A bound change that blocks one TDT may leave another violation feasible, violate a previously satisfied property, or compromise the original functional behavior. In this paper, functional behavior means the untimed discrete behavior characterized by functional equivalence. The central difficulty is not merely to block one counterexample, but to accept a repair only when it satisfies the full property set and preserves the original functional behavior. The running example in Section II illustrates this effect: repairing an early gear-engagement violation can make an acknowledgement deadline fail. A practical repair method should therefore generate small and inspectable candidate edits, but accept them only after full-property regression verification and an admissibility check based on untimed functional equivalence [4].

Existing repair techniques provide an important foundation. TARTAR [4] introduced TDT-based clock-bound repair by encoding a TDT into linear real arithmetic constraints, solving a MaxSMT problem for small bound changes, and checking admissibility by functional equivalence. Related work studies parametric timed automata, clock-guard repair, robustness analysis, and SMT-based repair [12]–[15]. However, these workflows are primarily trace-local and single-property in style. When repairs are generated property by property, the coupling among properties, counterexamples, and repairable bounds is not used to guide candidate generation, and a local edit may be rejected only after full validation.

This paper presents MPTA-Repair, an automated framework for simultaneous multi-property clock-bound repair of NTA models via iterative solving. The key idea is to treat each violated property and its TDT as a property-trace constraint, maintain a shared constraint pool, and use relations among properties, counterexamples, and repairable bounds to guide MaxSMT-based candidate generation. Each candidate is validated by full-property regression verification and an admissibility check based on functional equivalence. If validation

The implementation and datasets are publicly available at <https://github.com/zhougt/MPTA-Repair>.

exposes new violations, their TDTs are added to the pool and the solver is invoked again.

We evaluate MPTA-Repair on benchmarks derived from UPPAAL [16] and the XtaBenchmarkSuite [17], together with an industrial dataset constructed in this work. Faulty NTA instances are produced by controlled clock-bound mutations. Compared with an adapted iterative TARTAR-style baseline, MPTA-Repair improves the success rate from 72% to 93% on the benchmark dataset and from 20% to 76% on the industrial dataset.

This paper makes three contributions. First, it formulates simultaneous multi-property clock-bound repair for NTA models under full-property satisfaction and functional preservation. Second, it introduces an iterative MaxSMT workflow based on a shared property-trace constraint pool and relation-guided candidate generation. Third, it evaluates the framework on benchmark and industrial faulty models, showing that joint multi-property repair is more effective than an iterative trace-local baseline in this setting.

The remainder of the paper is organized as follows. Section II introduces the preliminaries and problem setting. Section III presents MPTA-Repair. Section IV reports the experimental evaluation. Section V concludes the paper.

II. PRELIMINARIES

A. Networks of Timed Automata and Properties

The paper follows standard timed-automata semantics [7]. This section establishes the notation used for multi-property clock-bound repair.

Definition 1 (Timed automaton). A timed automaton (TA) is a tuple

$$T = (L, \ell_0, C, \Sigma, \Theta, I), \quad (1)$$

where L is a finite set of locations, ℓ_0 is the initial location, C is a finite set of clocks, Σ is a set of action labels, Θ is a finite set of transitions, and I maps each location to a clock invariant. A clock constraint is a conjunction of atoms $x \sim b$, where $x \in C$, $\sim \in \{<, \leq, =, \geq, >\}$, and $b \in \mathbb{R}_{\geq 0}$. A transition $\theta = (\ell, g, a, R, \ell')$ moves from ℓ to ℓ' when guard g is satisfied, emits action a , and resets clocks in $R \subseteq C$. Each numerical constant b occurring in a guard or invariant is a clock-bound occurrence. MPTA-Repair only changes repairable clock-bound occurrences; it does not modify locations, transitions, synchronization labels, clock references, resets, data updates, or properties.

Definition 2 (NTA with property set). An NTA specification consists of a network of timed automata and the properties that the network must satisfy:

$$\begin{aligned} \mathcal{M} &= (\mathcal{N}, \Phi), \\ \mathcal{N} &= T_1 \parallel \dots \parallel T_k, \\ \Phi &= \{\varphi_1, \dots, \varphi_m\}. \end{aligned} \quad (2)$$

The network $\mathcal{N}$ models a real-time system by composing interacting timed automata, with synchronization semantics supplied by the target modeling language, such as UPPAAL [9]. The property set Φ is part of the system specification: it states

the timing and behavioral requirements that the NTA must satisfy. For convenience, $\mathcal{N} \models \Phi$ denotes that every property in Φ holds over the reachable states of $\mathcal{N}$.

The repair problem considered here is restricted to timed safety properties, which require a state predicate to hold over all reachable states. In UPPAAL-style notation, such properties often have the invariant form $\mathbb{A}[] p$, where p may combine location predicates, data predicates, and clock constraints. Clocks are central to timed safety properties because deadlines, timeouts, durations, and earliest or latest event times are expressed by numerical bounds over clock values.

B. Timed Diagnostic Traces and Accepted Repairs

Definition 3 (Timed diagnostic trace). Given an NTA specification $\mathcal{M} = (\mathcal{N}, \Phi)$ and a property $\varphi \in \Phi$, a timed diagnostic trace (TDT) for φ is a feasible finite timed execution of $\mathcal{N}$ that reaches a state violating φ . In practice, a TDT is returned by a model checker when verification of φ fails. It records the relevant locations and transitions, guards and invariants, clock resets, timing information such as delays or zones, and the violation point.

In MPTA-Repair, each TDT is treated as a property-trace constraint: the violated property identifies the requirement to be restored, and the trace identifies the timed behavior that makes the violation feasible. Section III uses such constraints from multiple properties in a joint solving process.

Definition 4 (Untimed functional equivalence). For an NTA $\mathcal{N}$, let $L_\mu(\mathcal{N})$ denote the untimed language obtained by abstracting from delays while preserving the discrete action or transition behavior feasible under the timing constraints. Two NTA models $\mathcal{N}_1$ and $\mathcal{N}_2$ are untimed functionally equivalent if

$$L_\mu(\mathcal{N}_1) = L_\mu(\mathcal{N}_2). \quad (3)$$

This notion follows the functional-equivalence criterion used for clock-bound repair [4]: a timing repair should not remove intended discrete behavior or add new discrete behavior, even though it changes timing constraints to eliminate violations.

Definition 5 (Accepted repair). Let $\mathcal{M}_{bad} = (\mathcal{N}_{bad}, \Phi)$ be a faulty NTA specification. A candidate repaired model $\mathcal{M}_{rep} = (\mathcal{N}_{rep}, \Phi)$ is obtained by changing only repairable clock-bound occurrences in guards and invariants. The candidate is accepted as a repair only if

$$\mathcal{N}_{rep} \models \bigwedge_{\varphi \in \Phi} \varphi \quad \wedge \quad L_\mu(\mathcal{N}_{bad}) = L_\mu(\mathcal{N}_{rep}). \quad (4)$$

The first conjunct corresponds to full-property regression verification: the repaired model must satisfy the entire property set, not only the property that produced the current TDT. The second conjunct is the admissibility check based on untimed functional equivalence; it prevents repairs that make properties true by disabling required functional behavior. MPTA-Repair uses this accepted-repair condition as the global validation criterion for candidates generated from local TDT evidence.

III. APPROACH

A. Overview

MPTA-Repair takes an NTA model M and a property set $\Phi = \{\varphi_1, \dots, \varphi_m\}$ as input. Instead of requiring manual variable specification, the framework extracts repairable bounds from numerical clock constraints in transition guards and location invariants. The final output is either a repaired model M' accompanied by concrete clock-bound edits, or a failure report if no valid full-property repair can be synthesized within the resource budget.

The framework operates on a local-generation and global-validation principle. While each TDT drives candidate generation, validation remains global to ensure that fixing one violation does not introduce regressions elsewhere. To leverage the interdependence where multiple violations share identical timing structures, MPTA-Repair maintains a joint pool of property-trace constraints. This joint structure enables the framework to exploit relations among properties, counterexamples, and repairable bounds to optimize the search space and guide the MaxSMT solver toward minimal parameter assignments.

Operationally, each refinement round begins by verifying M against Φ to collect TDTs for any violations. If violations are detected, the framework identifies active repairable bounds and constructs a joint MaxSMT problem from the active constraint pool. The synthesized parameter assignment is instantiated into a candidate model, which undergoes global validation via full-property regression verification and an admissibility check based on functional equivalence. A rejected candidate contributes newly exposed TDTs back into the pool to further constrain subsequent rounds, whereas a candidate passing both checks is returned as the valid repair M' . Figure 1 summarizes this counterexample-guided refinement loop from verification and relation-guided MaxSMT solving to global validation.

B. Joint Trace Encoding and Optimization Objective

MPTA-Repair follows the clock-bound variation strategy used in TARTAR [4]. Let S denote the set of repairable clock-bound occurrences. For each occurrence $s \in S$, the original numerical bound b_s in a transition guard or location invariant is associated with a variation variable δ_s . This ensures that the repaired constraint retains the original clock reference and comparison operator while incorporating a shifted value:

$$b'_s = b_s + \delta_s \quad (5)$$

The repair space is therefore strictly restricted to clock-bound constants, keeping transition structures and discrete updates unaltered.

MPTA-Repair builds on the trace-encoding formulation of TARTAR and extends it to the multi-property setting [4]. Given a TDT

$$\tau = \ell_0 \xrightarrow{d_0, e_0} \ell_1 \cdots \xrightarrow{d_{n-1}, e_{n-1}} \ell_n \quad (6)$$

a trace path is encoded as a conjunction of linear arithmetic constraints:

$$\begin{aligned} \text{Path}(\tau, \rho, \mathbf{z}_\tau) = & \text{Init}(\nu_0) \\ & \wedge \bigwedge_{i=0}^{n-1} \left(d_i \geq 0 \wedge \text{Inv}(\ell_i, \nu_i + d_i, \rho) \right. \\ & \wedge \text{Guard}(e_i, \nu_i + d_i, \rho) \\ & \left. \wedge \nu_{i+1} = \text{Reset}(\nu_i + d_i, R_i) \right) \end{aligned} \quad (7)$$

Here, ρ denotes the modified bounds, and $\mathbf{z}_\tau$ collects the delays and clock valuations. The reset equation $\nu_{i+1} = \text{Reset}(\nu_i + d_i, R_i)$ updates the valuation after transition e_i by resetting clocks in R_i and preserving the others. The original violation encoding combines path executability with the final property violation:

$$\text{Enc}(\tau, \varphi, \rho) = \text{Path}(\tau, \rho, \mathbf{z}_\tau) \wedge \text{Bad}_\varphi(\ell_n, \nu_n) \quad (8)$$

The predicate $\text{Bad}_\varphi(\ell_n, \nu_n)$ characterizes the violation state of φ at the final location. To block this violating behavior while preserving executability of the discrete path skeleton, the single-trace repair constraint uses two copies of the trace variables:

$$\begin{aligned} \Psi_{\tau, \varphi}^{\text{repair}}(\rho) = & \exists \mathbf{z}_\tau^+. \text{Path}(\tau, \rho, \mathbf{z}_\tau^+) \\ & \wedge \neg \exists \mathbf{z}_\tau^-. \left(\text{Path}(\tau, \rho, \mathbf{z}_\tau^-) \right. \\ & \left. \wedge \text{Bad}_\varphi(\ell_n^-, \nu_n^-) \right) \end{aligned} \quad (9)$$

The first part requires that the original discrete path skeleton still has at least one timed realization after repair. The second part requires that no assignment of delays and clock valuations can make that same path violate φ .

We extend this single-trace formulation to a joint, multi-property incremental setting. At refinement round k , counterexamples from all currently violated properties are accumulated into a trace pool

$$\mathcal{T}_k = \{(\tau_j, \varphi_j)\}_{j=1}^{m_k} \quad (10)$$

MPTA-Repair assembles a joint MaxSMT problem with the constraint system defined as

$$\begin{aligned} \Phi_k(\rho) = & \text{Dom}(\rho) \wedge \text{Block}_k(\rho) \\ & \wedge \bigwedge_{(\tau, \varphi) \in \text{Select}(\mathcal{T}_k)} \Psi_{\tau, \varphi}^{\text{repair}}(\rho) \end{aligned} \quad (11)$$

Here, $\text{Dom}(\rho)$ enforces static interval consistency and domain limits, such as $b'_s \geq 0$, and $\text{Select}(\mathcal{T}_k)$ represents a reduced subset of traces obtained after relation-aware filtering. When an instantiated candidate model fails verification, the constraint system is updated incrementally. The framework expands the pool via

$$\mathcal{T}_{k+1} = \mathcal{T}_k \cup \text{NewTDTs}(M(\rho_k), \Phi) \quad (12)$$

and updates the search-space block to exclude the failed assignment ρ_k over the set of repairable bounds S :

$$\text{Block}_{k+1}(\rho) = \text{Block}_k(\rho) \wedge \left(\bigvee_{s \in S} \delta_s \neq \delta_s^k \right) \quad (13)$$

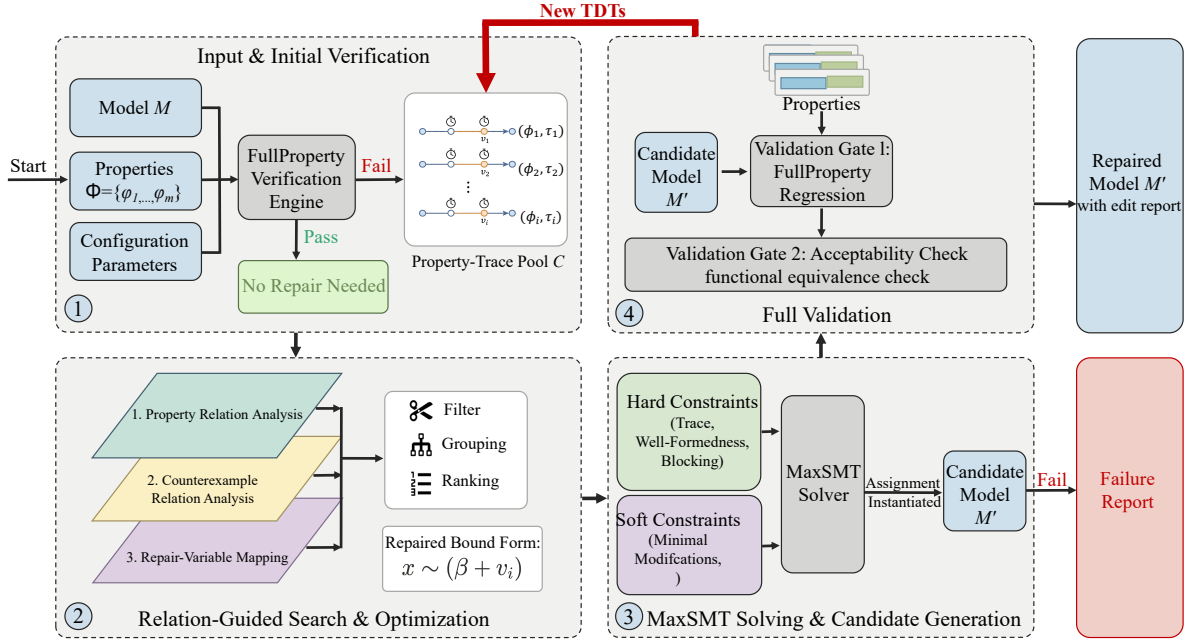


Fig. 1. Overview of the MPTA-Repair workflow.

where δ_s^k is the valuation assigned to variation variable δ_s in round k .

The optimization objective is governed by a hierarchical set of soft constraints. MPTA-Repair uses a lexicographic objective that first favors small repairs and then incorporates risk- and coverage-based priorities:

$$\text{lex min}_{\rho} \left(\begin{array}{l} \sum_{s \in S} \mathbf{1}\{\delta_s \neq 0\}, \\ \sum_{s \in S} |\delta_s|, \\ \sum_{s \in S} \text{risk}_s \cdot \mathbf{1}\{\delta_s \neq 0\}, \\ - \sum_{s \in S} \text{benefit}_s \cdot \mathbf{1}\{\delta_s \neq 0\} \end{array} \right) \quad (14)$$

Here, $\mathbf{1}\{\cdot\}$ denotes the indicator function, yielding 1 when the inner condition holds and 0 otherwise. The primary and secondary objectives minimize the number and total magnitude of altered clock bounds to produce concise repairs. The third term incorporates a penalty coefficient risk_s to discourage modifications prone to regression, while the final term uses a reward coefficient benefit_s to prioritize repair sites that cover a broader set of violated traces.

C. Relation-Guided Search Optimization

During refinement, accumulating violated properties, TDTs, and editable bounds expands the MaxSMT search space. To reduce redundant constraints, MPTA-Repair uses relation-guided optimization to filter, group, and rank constraints and repair variables. These heuristics only guide the search; correctness is still guaranteed through full-property regression verification and an admissibility check.

At the property layer, timed safety properties are normalized into an $A[]$ fragment over location predicates and clock or integer constraints. Equivalence is detected by the syntactic identity of normalized forms, while dominance is evaluated under logical subsumption, e.g., $A[] x \leq 5$ subsumes $A[] x \leq 10$ within the same structural context. MPTA-Repair then retains only the stronger properties for solving to optimize efficiency.

At the counterexample layer, MPTA-Repair maintains a TDT pool indexed by violated properties. Exact duplicates are removed, and the remaining traces are annotated with traversed transitions, visited locations, active clocks, and guard or invariant bounds. Traces with identical discrete transition sequences or overlapping clock zones are grouped to reveal common bottlenecks. A trace constraint is omitted from the active encoding only when it is an exact duplicate or soundly subsumed by an encoded trace constraint.

At the repair-bound layer, MPTA-Repair maps each active property-trace constraint to the clock bounds occurring in or affecting its encoding. This mapping restricts active MaxSMT decision variables to bounds relevant to the current counterexample pool and ranks them for candidate generation. Bounds with higher relational scores are prioritized because their modifications are more likely to satisfy the joint constraint system.

D. Validation, Admissibility, and Refinement

In MPTA-Repair, candidate validation determines whether a generated edit set can be accepted as a repair. After a candidate edit set is generated, the repaired model is verified against the entire property set Φ . If a candidate resolves the observed violation but induces new violations against the specification, it is rejected, indicating that the current property-trace pool

TABLE I
STATISTICS OF THE BENCHMARK AND INDUSTRIAL DATASETS.

Dataset	Original models	Mutated models	Avg. properties/model
Benchmark	9	54	4.1
Industrial	5	105	7.3

is insufficient. MPTA-Repair then extracts the newly observed TDTs and appends them to the pool, guiding the subsequent refinement round.

Candidates that pass full-property regression verification are further evaluated against an admissibility check based on functional equivalence [4]. This step ensures that clock-bound modifications do not satisfy properties in ways that fail to preserve the original functional behavior. Rather than enforcing timed-language equivalence, which would conflict with the goal of eliminating violating timed paths, the check is performed by constructing untimed automata from both the original and repaired models and comparing their accepted untimed languages [18]. A candidate is rejected if it adds or removes discrete functional paths, ensuring that only timing characteristics are modified.

The refinement loop therefore combines local candidate generation with global validation. Local MaxSMT formulations propose candidate assignments, full-property regression verification ensures complete specification compliance, and the admissibility check preserves the original functional behavior.

IV. EXPERIMENTAL EVALUATION

This section evaluates MPTA-Repair on both benchmark and industrial case studies. The evaluation follows the clock-bound repair setting defined in Section III, where only numerical bounds in guards and invariants are repairable. For comparison, we evaluate MPTA-Repair against an adapted iterative TARTAR-style baseline [4].

A. Evaluation Strategy and Datasets

We evaluate MPTA-Repair on two datasets, summarized in Table I. The benchmark dataset includes nine benchmark models derived from UPPAAL benchmark [16] and the XtaBenchmarkSuite [17], such as Elevator and FDDI. The industrial dataset constructed in this work contains five closely coupled NTA case-study models: Gear Control System for automotive transmission [19], SAE RoR for autonomous-driving decisions [20], Train Controller Alarm for railway emergency timing [21], Train Gate Controller for shared-resource concurrency [6], and Pacemaker for medical pacing [8]. Compared with the benchmark models, these case studies involve more tightly coupled timing constraints across multiple properties.

Since real-world NTA models with labeled clock-bound faults are difficult to obtain, we synthesize faulty instances using mutation [22], retaining only mutants that are syntactically valid, violate at least one timing safety property, and preserve the original functional behavior. Each repair instance is executed with a 10-minute timeout for both MPTA-Repair and the baseline.

TABLE II
REPAIR RESULTS ACROSS BENCHMARK AND INDUSTRIAL DATASETS.

Dataset	Method	Mutated models	Repaired	Success rate	Avg. time
Benchmark	Iterative baseline	54	39	72%	60.61s
	MPTA-Repair	54	50	93%	15.46s
Industrial	Iterative baseline	105	21	20%	35.59s
	MPTA-Repair	105	80	76%	85.28s

We compare MPTA-Repair against an adapted iterative TARTAR-style baseline [4], which was originally designed for single-property repair and operates via trace-local repair generation. This comparison allows us to examine whether using a joint counterexample pool improves repair robustness over localized trace analysis.

B. Experimental Results

The experimental evaluation demonstrates that MPTA-Repair outperforms the iterative baseline, particularly on complex industrial models. As shown in Table II, the baseline achieves a 72% success rate on the benchmark dataset, whereas MPTA-Repair reaches 93%. This performance gap widens on the industrial dataset, where the baseline repairs only 20% of the cases compared to 76% for MPTA-Repair. The baseline reports a lower average duration on the industrial dataset because unsuccessful runs typically terminate early when no valid candidate can be found.

The baseline remains competitive on simpler benchmarks, where injected timing faults are often localized to a small number of guards or invariants. However, its performance drops significantly on the industrial dataset, where, under concurrent component interactions and tightly coupled timing constraints, trace-local repairs often produce candidates that fail full-property regression verification or the admissibility check.

MPTA-Repair achieves higher success rates by integrating constraints from TDTs and properties into a joint generation process. Consequently, the framework produces coordinated modifications that pass both full-property regression verification and the admissibility check.

C. Discussion and Practical Lessons

The evaluation highlights two primary lessons for applying automated multi-property repair to NTA. First, full-property regression verification is necessary because industrial properties are often interdependent; a localized repair for one response-time property can inadvertently violate another safety property, such as maximum waiting bounds. MPTA-Repair avoids these regressions by requiring each candidate to satisfy the full property set. Second, the admissibility check is critical to distinguish valid repairs from behavior-breaking edits that satisfy timing properties only by removing intended functional behavior, such as preventing critical system actions. The admissibility check ensures that property satisfaction is not achieved at the expense of the original functional behavior.

Nevertheless, the framework can still fail to generate repairs in certain cases when it fails to find an admissible candidate within the time budget. This limitation is prominent in complex models like the Pacemaker case study, where tightly coupled timing constraints expand the search space, causing the solver to time out before finding a valid candidate.

V. CONCLUSION

This paper presented MPTA-Repair, an automated workflow for multi-property and multi-counterexample clock-bound repair of industrial NTA models. To overcome the limitations of isolated single-trace fixes, the method accumulates TDTs and exploits relations among properties, counterexamples, and repairable bounds to guide MaxSMT-based search. Each candidate is validated through full-property regression verification and an admissibility check based on functional equivalence [4], ensuring that accepted repairs preserve the original functional behavior.

Evaluations across benchmark and industrial datasets show that MPTA-Repair outperforms the iterative baseline, particularly in industrial systems with dense property interactions. The failure analysis further indicates that unsuccessful repairs mainly arise when the framework cannot find an admissible candidate within the time budget, especially as dense timing interactions enlarge the clock-bound repair search space.

Future work will primarily focus on extending the automated repair space beyond numerical clock bounds to include broader structural modifications, such as altering synchronization labels and redesigning discrete control logic. We also plan to integrate the admissibility check earlier in repair generation, so that functionally invalid candidates can be pruned before costly validation and large-scale industrial repair can be made more efficient.

REFERENCES

- [1] D. Basile, A. Fantechi, and I. Rosadi, “Formal analysis of the UNISIG safety application intermediate sub-layer: Applying formal methods to railway standard interfaces,” in *Proceedings of Formal Methods for Industrial Critical Systems*, ser. Lecture Notes in Computer Science, vol. 12863. Springer, 2021, pp. 174–190.
- [2] H. Lönn and P. Pettersson, “Formal verification of a TDMA protocol start-up mechanism,” in *Proceedings of the 1997 IEEE Pacific Rim International Symposium on Fault-Tolerant Systems*, 1997, pp. 235–242.
- [3] M. Mikučionis, K. G. Larsen, J. I. Rasmussen, B. Nielsen, A. Skou, S. U. Palm, J. S. Pedersen, and P. Hougard, “Schedulability analysis using Uppaal: Herschel-planck case study,” in *Proceedings of Leveraging Applications of Formal Methods, Verification and Validation*, ser. Lecture Notes in Computer Science, vol. 6416. Springer, 2010, pp. 175–190.
- [4] M. Kölbl, S. Leue, and T. Wies, “Clock bound repair for timed systems,” in *International Conference on Computer Aided Verification*. Springer, 2019, pp. 79–96.
- [5] T. Vogel, M. Carwehl, G. N. Rodrigues, and L. Grunske, “A property specification pattern catalog for real-time system verification with UPPAAL,” 2022, arXiv:2211.03817.
- [6] R. Alur and D. L. Dill, “A theory of timed automata,” *Theoretical Computer Science*, vol. 126, no. 2, pp. 183–235, 1994.
- [7] J. Bengtsson and W. Yi, “Timed automata: Semantics, algorithms and tools,” in *Lectures on Concurrency and Petri Nets*, ser. Lecture Notes in Computer Science. Springer, 2004, vol. 3098, pp. 87–124.
- [8] Z. Jiang, M. Pajic, A. Connolly, S. Dixit, and R. Mangharam, “Real-time heart model for implantable cardiac device validation and verification,” in *Proceedings of the Euromicro Conference on Real-Time Systems*. IEEE Computer Society, 2010, pp. 239–248.
- [9] K. G. Larsen, P. Pettersson, and W. Yi, “UPPAAL in a nutshell,” *International Journal on Software Tools for Technology Transfer*, vol. 1, no. 1–2, pp. 134–152, 1997.
- [10] C. Daws, A. Olivero, S. Tripakis, and S. Yovine, “The tool KRONOS,” in *Hybrid Systems III*, ser. Lecture Notes in Computer Science. Springer, 1996, vol. 1066, pp. 208–219.
- [11] A. E. Dalsgaard, R. R. Hansen, K. Y. Jørgensen, K. G. Larsen, M. C. Olesen, P. Olsen, and J. Srba, “opaal: A lattice model checker,” in *NASA Formal Methods*, ser. Lecture Notes in Computer Science, vol. 6617. Springer, 2011, pp. 487–493.
- [12] R. Alur, T. A. Henzinger, and M. Y. Vardi, “Parametric real-time reasoning,” in *Proceedings of the Twenty-Fifth Annual ACM Symposium on Theory of Computing*. ACM, 1993, pp. 592–601.
- [13] Étienne André, P. Arcaini, A. Gargantini, and M. Radavelli, “Repairing timed automata clock guards through abstraction and testing,” in *Proceedings of Tests and Proofs*, ser. Lecture Notes in Computer Science, vol. 11823. Springer, 2019, pp. 129–146.
- [14] J. Bendík, A. Sencan, E. A. Gol, and I. Černá, “Timed automata robustness analysis via model checking,” *Logical Methods in Computer Science*, vol. 18, no. 3, pp. 12:1–12:32, 2022.
- [15] M. Jose and R. Majumdar, “Bug-assist: Assisting fault localization in ANSI-C programs,” in *Proceedings of Computer Aided Verification*, ser. Lecture Notes in Computer Science, vol. 6806. Springer, 2011, pp. 504–509.
- [16] G. Behrmann, A. David, and K. G. Larsen, “A tutorial on Uppaal,” in *Formal Methods for the Design of Real-Time Systems*, ser. Lecture Notes in Computer Science, vol. 3185. Springer, 2004, pp. 200–236.
- [17] R. Farkas and G. Bergmann, “Towards reliable benchmarks of timed automata,” in *Proceedings of the 25th PhD Mini-Symposium*, 2018, pp. 20–23.
- [18] J. E. Hopcroft, R. Motwani, and J. D. Ullman, *Introduction to Automata Theory, Languages, and Computation*, 2nd ed. Addison-Wesley, 2001.
- [19] M. Lindahl, P. Pettersson, and W. Yi, “Formal design and analysis of a gear controller,” in *Proceedings of Tools and Algorithms for the Construction and Analysis of Systems*, ser. Lecture Notes in Computer Science, vol. 1384. Springer, 1998, pp. 281–297.
- [20] G. V. Alves, L. A. Dennis, and M. Fisher, “A double-level model checking approach for an agent-based autonomous vehicle and road junction regulations,” *Journal of Sensor and Actuator Networks*, vol. 10, no. 3, p. 41, 2021.
- [21] A. González, M. Cristiá, and C. Luna, “Verification of quantitative temporal properties in RealTime-DEVS,” *SIMULATION: Transactions of The Society for Modeling and Simulation International*, vol. 101, no. 10, pp. 1057–1076, 2025.
- [22] Y. Jia and M. Harman, “An analysis and survey of the development of mutation testing,” *IEEE Transactions on Software Engineering*, vol. 37, no. 5, pp. 649–678, 2011.