

MPTA-Repair: Simultaneous Multi-Property Clock Bound Repair for Networks of Timed Automata via Iterative Space Optimization

Guangtong Zhou, Run Chang, Ji Wu*

School of Computer Science and Engineering, Beihang University, Beijing, China

Email: zhougt@buaa.edu.cn, handsomeruncr@gmail.com, wuji@buaa.edu.cn

Abstract—Networks of timed automata (NTA) are widely used to model and verify real-time industrial systems, which are typically required to satisfy multiple properties simultaneously. When verification fails, repairing an NTA is challenging because a valid repair must modify erroneous clock constraints, satisfy all properties, and preserve the original functional behavior. We present MPTA-Repair, an automated repair framework for multi-property NTA. It uses an iterative, counterexample-guided MaxSMT approach to generate minimally modified candidates, exploiting relations among properties, counterexamples, and repair variables to optimize search. Each candidate is validated by full-property regression verification and an admissibility check based on functional equivalence. Compared to the TARTAR baseline, MPTA-Repair achieves higher effectiveness on faulty models generated via our proposed mutation strategy. Specifically, it improves the success rate from 72% to 93% on benchmarks derived from UPPAAL and the XtaBenchmarkSuite, and from 20% to 76% on an industrial dataset constructed in this work.

Index Terms—networks of timed automata, automated model repair, MaxSMT, multi-property verification

I. INTRODUCTION

Real-time systems are widely used in domains such as railway control, communication protocols, and aerospace scheduling [1]–[3]. Their correctness depends not only on whether they exhibit the intended discrete behavior, but also on whether relevant actions occur within the required time bounds. Timed automata (TAs) provide a formalism for modeling these two aspects by extending finite-state automata with real-valued clocks and timing constraints that govern both the occurrence of transitions and the time spent in locations [4], [5]. Systems composed of multiple interacting components are commonly represented as networks of timed automata (NTA). In practice, NTA are typically verified against sets of properties that collectively specify the required system behavior, including safety and liveness requirements [6], [7].

Among the properties checked for an NTA, this work focuses on timed safety properties. This scope is motivated by the diagnostic basis of clock-bound repair: violations of such properties can be witnessed by finite timed executions, which provide concrete evidence for subsequent repair analysis [7], [8]. Timed safety properties specify bounded timing requirements, such as deadlines, limits on waiting times, and

valid execution durations. Guaranteed timed safety presents a significant challenge in complex NTAs with numerous interacting components. In these models, multiple real-valued clocks advance simultaneously, and their constraints determine when transitions may occur and how long the system may remain in a location. As a result, an incorrect numerical clock bound may affect executions across multiple components, causing an action to occur too early, too late, or for an invalid duration, and thereby violating a timed safety property [2], [8], [9]. When such a violation is detected, the erroneous clock bounds must be repaired while preserving the intended functional behavior of the system.

When a property violation is detected, real-time model checkers such as UPPAAL, Kronos, and opaal can return a timed counterexample, commonly called a timed diagnostic trace (TDT) [10]–[12]. A TDT describes a timed sequence of transitions and delays that leads the NTA from its initial state to a state in which the property is violated. It shows how the violation occurs, but does not explain which clock bounds caused it or how those bounds should be changed. This repair problem becomes more difficult when an NTA must satisfy multiple properties simultaneously. Even if a candidate clock-bound change resolves the violation witnessed by one TDT, other violations may remain, and the same change may cause a previously satisfied property to fail or compromise the original functional behavior. Therefore, a candidate repair cannot be accepted solely because it resolves an individual violation; it must also be validated against the full property set and checked for admissibility based on functional equivalence.

Parameterization and constraint solving provide a general basis for repairing timing constraints in timed automata. A common strategy is to replace selected clock bounds with parameters and synthesize valuations under which the resulting model satisfies a specified property. Knapik and Penczek encode bounded executions of parametric timed automata as SMT formulas and synthesize parameter valuations for reachability properties [13]. TARTAR [8] encodes the timing constraints of a TDT in linear real arithmetic and uses MaxSMT to generate candidate repairs with minimal changes to clock bounds. It then performs an admissibility check based on functional equivalence to determine whether a candidate preserves the original functional behavior. Although these methods show that parameterization and constraint solving

The implementation and datasets are publicly available at <https://github.com/zhougt/MPTA-Repair>.

can support automated model repair, their repair or synthesis procedures are formulated for an individual property or TDT rather than a property set that must be satisfied simultaneously. As a result, they do not provide an iterative repair strategy or search optimization specifically designed for the multi-property setting.

To address this limitation, we present MPTA-Repair, an automated framework for simultaneous multi-property clock-bound repair of NTAs. The main contributions of this work are as follows. First, we formalize simultaneous multi-property clock-bound repair for NTA models, together with the validity conditions of full-property satisfaction and preservation of the original functional behavior. Second, we propose an iterative, counterexample-guided MaxSMT framework that maintains a joint property-trace constraint pool and exploits relations among properties, counterexamples, and repair variables to optimize the repair search. Each candidate repair model undergoes full-property regression verification and an admissibility check based on functional equivalence. If regression verification reveals a new violation, the corresponding TDT is added to the property-trace constraint pool, and the updated pool is used in the next repair iteration. Third, we propose a mutation strategy for constructing faulty NTA models and evaluate MPTA-Repair against the TARTAR baseline [8] on benchmark models from UPPAAL examples [14] and the XtaBenchmark-Suite [15], as well as industrial models constructed in this work, showing higher repair success rates in both settings.

The remainder of the paper is organized as follows. Section II introduces the notation and accepted-repair condition. Section III presents MPTA-Repair. Section IV reports the experimental evaluation. Section V concludes the paper.

II. PRELIMINARIES

A. Networks of Timed Automata and Properties

The paper follows standard timed-automata semantics [5]. This section establishes the notation used for multi-property clock-bound repair.

Definition 1 (Timed automaton). A timed automaton (TA) is a tuple

$$T = (L, \ell_0, C, \Sigma, \Theta, I), \quad (1)$$

where L is a finite set of locations, ℓ_0 is the initial location, C is a finite set of clocks, Σ is a set of action labels, $\mathcal{B}(C)$ is the set of clock constraints over C , $\Theta \subseteq L \times \mathcal{B}(C) \times \Sigma \times 2^C \times L$ is a finite set of transitions, and $I : L \rightarrow \mathcal{B}(C)$ maps each location to a clock invariant. A clock constraint is a conjunction of atoms $x \sim b$, where $x \in C$, $\sim \in \{<, \leq, =, \geq, >\}$, and $b \in \mathbb{R}_{\geq 0}$. A transition $\theta = (\ell, g, a, R, \ell')$ moves from ℓ to ℓ' when guard g is satisfied, emits action a , and resets clocks in $R \subseteq C$. As mentioned above, this work focuses on repairing clock bounds. Specifically, a clock bound is the numerical constant b in an atomic clock constraint $x \sim b$ appearing in a transition guard or location invariant.

Definition 2 (NTA with property set). An NTA specification consists of a network of timed automata and the properties that the network must satisfy:

$$\begin{aligned} M &= (N, \Phi), \\ N &= T_1 \parallel \dots \parallel T_k, \\ \Phi &= \{\varphi_1, \dots, \varphi_m\}. \end{aligned} \quad (2)$$

The network N models a real-time system by composing interacting timed automata, with synchronization semantics supplied by the target modeling language, such as UPPAAL [10]. The property set Φ is part of the system specification: it states the timing and behavioral requirements that the NTA must satisfy. The notation $N \models \Phi$ means that N satisfies every property in Φ :

$$N \models \Phi \iff \forall \varphi_i \in \Phi, N \models \varphi_i. \quad (3)$$

The repair problem considered here is restricted to timed safety properties, which require a state predicate to hold in every reachable state. In UPPAAL-style notation, such a property has the invariant form $A[] p$, where p is a state predicate composed of location predicates and clock constraints. For example, $A[] (\text{Controller.Waiting} \implies x \leq 10)$ requires that clock x never exceed 10 while the controller is in the `Waiting` location.

B. Timed Diagnostic Traces

Definition 3 (Timed diagnostic trace). Given an NTA specification $M = (N, \Phi)$ and a property $\varphi \in \Phi$, a timed diagnostic trace (TDT) for φ is a feasible finite timed execution of N that reaches a state violating φ . When verification of φ fails, a real-time model checker can generate a TDT as a counterexample. In UPPAAL, such a counterexample consists of a sequence of states and transitions, together with clock constraints describing the associated timing information [10].

In MPTA-Repair, each TDT is treated as a property-trace constraint: the violated property identifies the requirement to be restored, and the trace identifies the timed behavior that makes the violation feasible. Section III uses such constraints from multiple properties in a joint solving process.

C. Functional Equivalence and Repair Acceptance

Following TARTAR [8], we adopt equality of untimed languages as the criterion for functional equivalence. Under this criterion, a timing repair should neither remove existing functional behavior of the original model nor introduce new functional behavior, even though it changes timing constraints to eliminate violations.

Definition 4 (Functional equivalence). For an NTA N , let $L_\mu(N)$ denote the untimed language obtained by abstracting from delays while preserving the discrete behavior feasible under the timing constraints. Two NTA models N_1 and N_2 are functionally equivalent if

$$L_\mu(N_1) = L_\mu(N_2). \quad (4)$$

Definition 5 (Accepted repair). Let $M_{bad} = (N_{bad}, \Phi)$ be a faulty NTA specification. A candidate repair $M_{rep} =$

(N_{rep}, Φ) is obtained by changing only repairable clock bounds in guards and invariants. The candidate is accepted if and only if

$$N_{rep} \models \Phi \quad \wedge \quad L_\mu(N_{bad}) = L_\mu(N_{rep}). \quad (5)$$

The first conjunct corresponds to full-property regression verification: the model must satisfy the entire property set, not only the property that produced the current TDT. The second conjunct is the admissibility check based on untimed functional equivalence; it prevents repairs that make properties true by disabling required functional behavior.

III. APPROACH

A. Overview

MPTA-Repair takes as input $M = (N, \Phi)$, where N is an NTA and $\Phi = \{\varphi_1, \dots, \varphi_m\}$ is a set of timed safety properties, and modifies only repairable numerical clock bounds in transition guards and location invariants. It returns $M' = (N', \Phi)$ that satisfies the accepted-repair condition in Eq. (5); otherwise, it reports that no accepted repair was found within the resource budget.

The properties in Φ may not be independent, as different properties may be affected by the same clock bounds. Their TDTs may be identical, share part of an execution path, or depend on the same repairable bounds. A repair derived from one TDT may leave other violations unresolved or cause a satisfied property to fail. Treating TDTs separately ignores these relations and may lead to repeated, conflicting repairs.




Fig. 1. Overview of the MPTA-Repair workflow.

Figure 1 summarizes the MPTA-Repair framework. Initial full-property verification either confirms that no repair is needed or initializes a joint property-trace constraint pool with TDTs. MPTA-Repair analyzes relations among properties, TDTs, and repairable bounds, removes redundant constraints, and solves constraints that share bounds jointly. The constraints are encoded as a MaxSMT problem, and the resulting

assignment is applied to N to construct a candidate model. MPTA-Repair extends TARTAR's single-TDT MaxSMT repair procedure [8] by jointly processing related constraints and using newly generated TDTs in subsequent repair iterations.

Each candidate undergoes full-property regression verification and the admissibility check based on functional equivalence. If regression verification reveals new violations, the corresponding TDTs are added to the joint pool. Any assignment that fails either validation step is excluded from subsequent iterations. The process repeats until an accepted repair is found or the time budget is exhausted. Figure 2 details the relation analysis, MaxSMT solving, and validation stages.

B. Joint Trace Encoding

MPTA-Repair uses the clock-bound variation representation introduced in TARTAR [8] as the basis for its joint multi-property encoding and iterative repair procedure. Let S denote the set of repairable clock bounds. For each clock bound $s \in S$, the original numerical bound b_s in a transition guard or location invariant is associated with a variation variable δ_s . The repaired numerical bound b'_s is defined as:

$$b'_s = b_s + \delta_s \quad (6)$$

MPTA-Repair builds on the trace-encoding formulation of TARTAR and extends it to the multi-property setting [8]. Consider a TDT with n discrete transitions:

$$\tau = \ell_0 \xrightarrow{d_0, \theta_0} \ell_1 \cdots \xrightarrow{d_{n-1}, \theta_{n-1}} \ell_n \quad (7)$$

Here, d_i is the delay before transition θ_i , where $\theta_i = (\ell_i, g_i, a_i, R_i, \ell_{i+1})$. Let $\rho = (b'_s)_{s \in S}$ denote the modified-bound vector, and let $\mathbf{z}_\tau = (d_0, \dots, d_{n-1}, \nu_0, \dots, \nu_n)$ collect the delay values and clock valuations along this execution, where ν_i is the clock valuation at location ℓ_i . Then the trace path is encoded as a conjunction of linear arithmetic constraints:

$$\begin{aligned} \text{Path}(\tau, \rho, \mathbf{z}_\tau) = & \text{Init}(\nu_0) \\ & \wedge \bigwedge_{i=0}^{n-1} \left(d_i \geq 0 \wedge \text{Inv}(\ell_i, \nu_i + d_i, \rho) \right. \\ & \wedge \text{Guard}(\theta_i, \nu_i + d_i, \rho) \\ & \left. \wedge \nu_{i+1} = \text{Reset}(\nu_i + d_i, R_i) \right) \end{aligned} \quad (8)$$

Here, $\text{Init}(\nu_0)$ constrains the initial clock valuation, Inv evaluates the invariant $I(\ell_i)$, and Guard evaluates the guard g_i of transition θ_i under the given valuation and modified bounds. The reset equation $\nu_{i+1} = \text{Reset}(\nu_i + d_i, R_i)$ updates the valuation after transition θ_i by resetting clocks in R_i and preserving the others. The predicate $\text{Bad}_\varphi(\ell_n, \nu_n)$ characterizes the violation state of φ at the final location. To block this violating behavior while preserving executability of the discrete path skeleton, the single-trace repair constraint uses

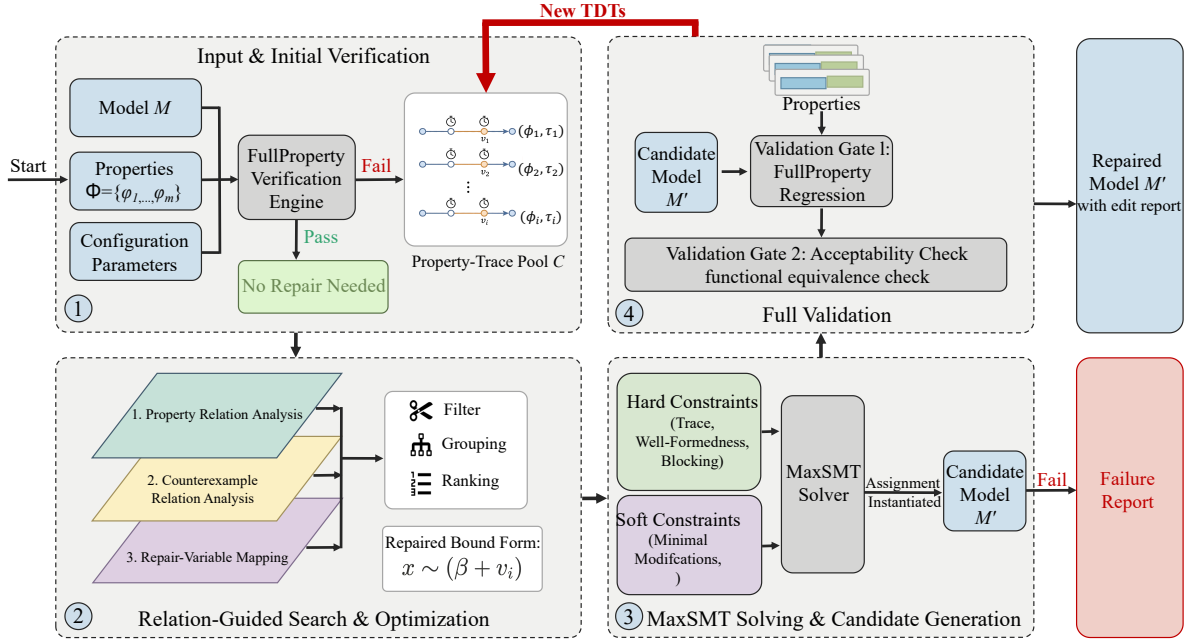


Fig. 2. Overview of the MPTA-Repair workflow.

two disjoint copies $\mathbf{z}_\tau^+$ and $\mathbf{z}_\tau^-$ of the delay and clock-valuation variables:

$$\begin{aligned} \Psi_{\tau, \varphi}^{\text{repair}}(\rho) = & \exists \mathbf{z}_\tau^+. \text{Path}(\tau, \rho, \mathbf{z}_\tau^+) \\ & \wedge \neg \exists \mathbf{z}_\tau^-. (\text{Path}(\tau, \rho, \mathbf{z}_\tau^-) \\ & \wedge \text{Bad}_\varphi(\ell_n, \nu_n^-)) \end{aligned} \quad (9)$$

The first part requires that the original discrete path skeleton still has at least one timed realization after repair. This preserves the observed discrete behavior: the repair must not remove the transition sequence simply to avoid the violation. The second part requires that no assignment of delays and clock valuations can make that same path violate φ .

We extend this single-trace formulation to a joint, multi-property incremental setting. At refinement round k , counterexamples from all currently violated properties are accumulated into a trace pool

$$\mathcal{T}_k = \{(\tau_j, \varphi_j)\}_{j=1}^{m_k} \quad (10)$$

where $m_k = |\mathcal{T}_k|$ is the number of property-trace pairs in the pool at round k . MPTA-Repair assembles a joint MaxSMT problem with the constraint system defined as

$$\begin{aligned} \mathcal{F}_k(\rho) = & \text{Dom}(\rho) \wedge \text{Block}_k(\rho) \\ & \wedge \bigwedge_{(\tau, \varphi) \in \text{Select}(\mathcal{T}_k)} \Psi_{\tau, \varphi}^{\text{repair}}(\rho) \end{aligned} \quad (11)$$

Here, $\text{Dom}(\rho)$ enforces static interval consistency and domain limits, such as $b'_s \geq 0$; $\text{Block}_k(\rho)$ collects the blocking clauses accumulated before round k and is initially true; and $\text{Select}(\mathcal{T}_k)$ represents a reduced subset of traces obtained after relation-aware filtering. When an instantiated candidate NTA

fails verification, the constraint system is updated incrementally. The framework expands the pool via

$$\mathcal{T}_{k+1} = \mathcal{T}_k \cup \text{NewTDTs}(N(\rho_k), \Phi) \quad (12)$$

where ρ_k is the modified-bound assignment produced in round k , $N(\rho_k)$ is the candidate network obtained by applying ρ_k to N , and $\text{NewTDTs}(N(\rho_k), \Phi)$ is the set of newly exposed property-trace pairs returned by full-property verification. The framework then updates the search-space block to exclude the failed assignment ρ_k over the set of repairable bounds S :

$$\text{Block}_{k+1}(\rho) = \text{Block}_k(\rho) \wedge \left(\bigvee_{s \in S} \delta_s \neq \delta_s^k \right) \quad (13)$$

where δ_s^k is the valuation assigned to variation variable δ_s in round k .

C. Relation-Guided Search Optimization

During refinement, accumulating violated properties, TDTs, and repairable bounds increases the size and complexity of the MaxSMT problem. To improve solving efficiency, MPTA-Repair uses relation-guided optimization to filter and group constraints, and to rank repairable bounds. These heuristics only guide the search; correctness is still guaranteed through full-property regression verification and an admissibility check.

MPTA-Repair analyzes relations at the property, counterexample, and repair-bound levels to reduce and organize the active MaxSMT problem. Timed safety properties are grouped according to the state predicates and clock constraints appearing in their $(\mathbb{A}[\cdot], p)$ queries; identical queries are treated as duplicates, while tighter bounds are given higher priority among properties with the same structure. Each TDT is stored

together with the property it violates, and property-trace pairs that share transitions or repair variables are grouped for joint solving. The framework then maps each active property-trace constraint to the repairable clock-bound occurrences in its encoding. A bound is given higher priority if it occurs in newly added TDTs or is associated with more TDTs and violated properties. This ranking guides active variable selection and candidate generation, whereas the number and magnitude of clock-bound changes are handled by the MaxSMT optimization objective. All properties remain in Φ for full-property regression verification, and a constraint is removed from the active encoding only when it is a duplicate or logically implied by another active constraint.

In iterative solving, MPTA-Repair caches trace-level quantifier elimination (QE). For each property-trace pair, it eliminates the delay and trace-local clock variables in $\Psi_{\tau, \varphi}^{\text{repair}}(\rho)$ and memoizes the resulting quantifier-free obligation using the trace, property, and repair-space configuration as the key. Later rounds reuse cached obligations for old TDTs and invoke QE only for newly discovered pairs; blocking clauses over ρ are added after QE and do not invalidate the cache.

The optimization objective is governed by a hierarchical set of soft constraints. MPTA-Repair uses a lexicographic objective that first favors small repairs and then incorporates risk- and coverage-based priorities:

$$\text{lex min}_{\rho} \begin{pmatrix} \sum_{s \in S} \mathbf{1}\{\delta_s \neq 0\}, \\ \sum_{s \in S} |\delta_s|, \\ \sum_{s \in S} \text{risk}_s \cdot \mathbf{1}\{\delta_s \neq 0\}, \\ - \sum_{s \in S} \text{benefit}_s \cdot \mathbf{1}\{\delta_s \neq 0\} \end{pmatrix} \quad (14)$$

Here, $\mathbf{1}\{\cdot\}$ is the indicator function. The first two objectives minimize the number and total magnitude of altered clock bounds. The last two terms penalize high-risk edits and prioritize bounds covering more violated traces.

IV. EXPERIMENTAL EVALUATION

This section evaluates MPTA-Repair on both benchmark and industrial case studies. The evaluation follows the clock-bound repair setting defined in Section III, where only numerical bounds in guards and invariants are repairable. For comparison, we evaluate MPTA-Repair against an adapted iterative TARTAR-style baseline [8].

A. Evaluation Setup and Research Questions

We evaluate MPTA-Repair on two datasets, summarized in Table I. The benchmark dataset includes nine benchmark models derived from UPPAAL benchmark [14] and the XtaBenchmarkSuite [15], such as Elevator and FDDI. The industrial dataset constructed in this work contains five closely coupled NTA case-study models: Gear Control System for automotive transmission [16], SAE RoR for autonomous-driving decisions

TABLE I
STATISTICS OF THE BENCHMARK AND INDUSTRIAL DATASETS.

Dataset	Original models	Mutated models	Avg. properties/model
Benchmark	9	54	4.1
Industrial	5	105	7.3

TABLE II
REPAIR RESULTS ACROSS BENCHMARK AND INDUSTRIAL DATASETS.

Dataset	Method	Mutated models	Repaired	Success rate	Avg. time
Benchmark	Iterative baseline	54	39	72%	60.61s
	MPTA-Repair	54	50	93%	15.46s
Industrial	Iterative baseline	105	21	20%	35.59s
	MPTA-Repair	105	80	76%	85.28s

[17], Train Controller Alarm for railway emergency timing [18], Train Gate Controller for shared-resource concurrency [4], and Pacemaker for medical pacing [9]. Compared with the benchmark models, these case studies involve more tightly coupled timing constraints across multiple properties.

Since real-world NTA models with labeled clock-bound faults are difficult to obtain, we synthesize faulty instances using mutation [19], retaining only mutants that are syntactically valid, violate at least one timing safety property, and preserve the original functional behavior. Each repair instance is executed with a 10-minute timeout for both MPTA-Repair and the baseline.

We compare MPTA-Repair against an adapted iterative TARTAR-style baseline [8], which was originally designed for single-property repair and operates via trace-local repair generation. The evaluation is organized around the following research questions:

RQ1. How effective is MPTA-Repair compared with the iterative TARTAR-style baseline?

RQ2. How does repair effectiveness change from benchmark models to industrial NTA models with more tightly coupled timing constraints?

RQ3. What practical lessons do the results reveal about full-property validation, admissibility checking, and remaining repair failures?

B. Results by Research Question

RQ1: Overall repair effectiveness. Table II shows that MPTA-Repair outperforms the iterative baseline on both datasets. On the benchmark dataset, the baseline repairs 39 of 54 faulty instances, achieving a 72% success rate, whereas MPTA-Repair repairs 50 instances and reaches 93%. This result indicates that joint property-trace candidate generation improves repair effectiveness even when injected timing faults are relatively localized.

RQ2: Effect on industrial NTA models. The advantage of MPTA-Repair becomes more pronounced on the industrial dataset. The iterative baseline repairs only 21 of 105 faulty instances, corresponding to a 20% success rate, while MPTA-Repair repairs 80 instances and reaches 76%. The baseline

reports a lower average duration on the industrial dataset because unsuccessful runs typically terminate early when no valid candidate can be found. Its success rate drops because trace-local repairs often produce candidates that fail full-property regression verification or the admissibility check under concurrent component interactions and tightly coupled timing constraints.

RQ3: Practical lessons and failure cases. MPTA-Repair achieves higher success rates by integrating constraints from TDTs and properties into a joint generation process. Consequently, the framework produces coordinated modifications that pass both full-property regression verification and the admissibility check. The results highlight two practical lessons for applying automated multi-property repair to NTA. First, full-property regression verification is necessary because industrial properties are often interdependent; a localized repair for one response-time property can inadvertently violate another safety property, such as maximum waiting bounds. MPTA-Repair avoids these regressions by requiring each candidate to satisfy the full property set. Second, the admissibility check is critical to distinguish valid repairs from behavior-breaking edits that satisfy timing properties only by removing intended functional behavior, such as preventing critical system actions.

Nevertheless, the framework can still fail to generate repairs in certain cases when it fails to find an admissible candidate within the time budget. This limitation is prominent in complex models like the Pacemaker case study, where tightly coupled timing constraints expand the search space, causing the solver to time out before finding a valid candidate.

V. CONCLUSION

This paper presented MPTA-Repair, an automated framework for simultaneous multi-property clock-bound repair of NTAs. It uses an iterative, counterexample-guided MaxSMT procedure that maintains a joint property-trace constraint pool and exploits relations among properties, counterexamples, and repair variables to optimize the repair search. Each candidate undergoes full-property regression verification and the admissibility check based on functional equivalence [8]; any newly revealed TDT is added to the pool for the next repair iteration.

The experimental results show that MPTA-Repair improves the repair success rate over the TARTAR baseline from 72% to 93% on the benchmark dataset and from 20% to 76% on the industrial dataset. The remaining failures mainly occur on complex models, where interactions among timing constraints enlarge the clock-bound repair space and prevent an admissible candidate from being found within the time budget.

Future work will improve repair efficiency by checking candidate admissibility earlier, allowing functionally invalid candidates to be discarded before full validation. We will also extend the repair space beyond numerical clock bounds to synchronization labels and discrete transition structures, together with new admissibility conditions for preserving the intended functional behavior.

REFERENCES

- [1] D. Basile, A. Fantechi, and I. Rosadi, “Formal analysis of the UNISIG safety application intermediate sub-layer: Applying formal methods to railway standard interfaces,” in *Proceedings of Formal Methods for Industrial Critical Systems*, ser. Lecture Notes in Computer Science, vol. 12863. Springer, 2021, pp. 174–190.
- [2] H. Lönn and P. Pettersson, “Formal verification of a TDMA protocol start-up mechanism,” in *Proceedings of the 1997 IEEE Pacific Rim International Symposium on Fault-Tolerant Systems*, 1997, pp. 235–242.
- [3] M. Mikučionis, K. G. Larsen, J. I. Rasmussen, B. Nielsen, A. Skou, S. U. Palm, J. S. Pedersen, and P. Hougard, “Schedulability analysis using Uppaal: Herschel-planck case study,” in *Proceedings of Leveraging Applications of Formal Methods, Verification and Validation*, ser. Lecture Notes in Computer Science, vol. 6416. Springer, 2010, pp. 175–190.
- [4] R. Alur and D. L. Dill, “A theory of timed automata,” *Theoretical Computer Science*, vol. 126, no. 2, pp. 183–235, 1994.
- [5] J. Bengtsson and W. Yi, “Timed automata: Semantics, algorithms and tools,” in *Lectures on Concurrency and Petri Nets*, ser. Lecture Notes in Computer Science. Springer, 2004, vol. 3098, pp. 87–124.
- [6] E. M. Clarke, T. A. Henzinger, H. Veith, and R. Bloem, Eds., *Handbook of Model Checking*. Springer, 2018.
- [7] T. Vogel, M. Carwehl, G. N. Rodrigues, and L. Grunske, “A property specification pattern catalog for real-time system verification with UPPAAL,” 2022, arXiv:2211.03817.
- [8] M. Kölbl, S. Leue, and T. Wies, “Clock bound repair for timed systems,” in *International Conference on Computer Aided Verification*. Springer, 2019, pp. 79–96.
- [9] Z. Jiang, M. Pajic, A. Connolly, S. Dixit, and R. Mangharam, “Real-time heart model for implantable cardiac device validation and verification,” in *Proceedings of the Euromicro Conference on Real-Time Systems*. IEEE Computer Society, 2010, pp. 239–248.
- [10] K. G. Larsen, P. Pettersson, and W. Yi, “UPPAAL in a nutshell,” *International Journal on Software Tools for Technology Transfer*, vol. 1, no. 1–2, pp. 134–152, 1997.
- [11] C. Daws, A. Olivero, S. Tripakis, and S. Yovine, “The tool KRONOS,” in *Hybrid Systems III*, ser. Lecture Notes in Computer Science. Springer, 1996, vol. 1066, pp. 208–219.
- [12] A. E. Dalsgaard, R. R. Hansen, K. Y. Jørgensen, K. G. Larsen, M. C. Olesen, P. Olsen, and J. Srba, “opaal: A lattice model checker,” in *NASA Formal Methods*, ser. Lecture Notes in Computer Science, vol. 6617. Springer, 2011, pp. 487–493.
- [13] M. Knapik and W. Penczek, “SMT-based parameter synthesis for parametric timed automata,” in *Challenging Problems and Solutions in Intelligent Systems*, ser. Studies in Computational Intelligence, G. de Tré, P. Grzegorzewski, J. Kacprzyk, J. W. Owsinski, W. Penczek, and S. Zadrozny, Eds. Cham: Springer, 2016, vol. 634, pp. 3–21.
- [14] G. Behrmann, A. David, and K. G. Larsen, “A tutorial on Uppaal,” in *Formal Methods for the Design of Real-Time Systems*, ser. Lecture Notes in Computer Science, vol. 3185. Springer, 2004, pp. 200–236.
- [15] R. Farkas and G. Bergmann, “Towards reliable benchmarks of timed automata,” in *Proceedings of the 25th PhD Mini-Symposium*, 2018, pp. 20–23.
- [16] M. Lindahl, P. Pettersson, and W. Yi, “Formal design and analysis of a gear controller,” in *Proceedings of Tools and Algorithms for the Construction and Analysis of Systems*, ser. Lecture Notes in Computer Science, vol. 1384. Springer, 1998, pp. 281–297.
- [17] G. V. Alves, L. A. Dennis, and M. Fisher, “A double-level model checking approach for an agent-based autonomous vehicle and road junction regulations,” *Journal of Sensor and Actuator Networks*, vol. 10, no. 3, p. 41, 2021.
- [18] A. González, M. Cristiá, and C. Luna, “Verification of quantitative temporal properties in RealTime-DEVS,” *SIMULATION: Transactions of The Society for Modeling and Simulation International*, vol. 101, no. 10, pp. 1057–1076, 2025.
- [19] Y. Jia and M. Harman, “An analysis and survey of the development of mutation testing,” *IEEE Transactions on Software Engineering*, vol. 37, no. 5, pp. 649–678, 2011.